

Engineering Excellence J1

DATA LAKE BENCHMARK TESTING

PREREQUISITES	2
CONFIGURATIONS.....	3
HDFS TESTS	4
1. TestDFSIO	4
2. TERA SORT BENCHMARK SUITE	6
3. NAMENODE BENCHMARK (NNBENCH).....	7
4. MAP REDUCE TESTS	11
• MapReduce Benchmark (MRBench)	11
CONCLUSION.....	12

Overview

This document provides overview of one of the Hadoop cluster evaluation process.

Hadoop is a framework which enables applications to work on large data on clusters which can have n numbers of nodes built of commodity hardware. It provides a distributed file system (HDFS) that stores data on the computed nodes, providing very high aggregate bandwidth across the cluster. In addition, Hadoop implements a parallel computational paradigm named MapReduce which divides the application into many small fragments of work, each of which may be executed or re executed on any node in the cluster.

The benchmarking of the cluster which means assessing its performances by running a variety of jobs each focused on a specific field (indexing, querying, predictive statistics, machine learning ...).

This benchmark measures response on a handful of relational queries: scans, aggregations, joins, and UDF's, across different data sizes. Keep in mind that these systems have very different sets of capabilities. MapReduce-like systems (Shark/Hive) target flexible and large-scale computation, supporting complex User Defined Functions (UDF's), tolerating failures, and scaling to thousands of nodes. Traditional MPP databases are strictly SQL compliant and heavily optimized for relational queries. The workload here is simply one set of queries that most of these systems these can complete.

The Hadoop distribution comes with a number of benchmarks, which are bundled in `hadoop-test.jar` and `hadoop-examples.jar`. The four benchmarks we will be looking at in more details are TestDFSIO, nnbench, mrbench (in `hadoop-test.jar`) and TeraGen / TeraSort / TeraValidate (in `hadoop-examples.jar`).

PREREQUISITES

Run the following command before executing benchmark tests.

1. Create directory `/benchmarks` in hdfs

```
sudo -u hdfs hdfs dfs -mkdir /benchmarks
```

2. Provide ownership on newly created directory to user who is going to execute benchmark tests.

```
sudo -u hdfs hdfs dfs -chown -R hdfs:hdfs /benchmarks
```

CONFIGURATIONS

Group	Property	Property Name	Property Value
HDFS Block size		dfs.block.size	134 MB
Java Heap Size of NameNode			1GB
Java Heap Size of DataNode			1GB
ResourceManager	Container Memory Minimum	yarn.scheduler.minimum-allocation-mb	512 MB
ResourceManager	Container Memory Maximum	yarn.scheduler.maximum-allocation-mb	1024MB
NodeManager	Container Memory	yarn.nodemanager.resource.memory-mb	4096MB
Gateway	Map Task Memory	mapreduce.map.memory.mb	
Gateway	Map Task Java Opts Base	mapreduce.map.java.opts	
Gateway	Reduce Task Memory	mapreduce.reduce.memory.mb	
Gateway	Reduce Task Java Opts Base	mapreduce.reduce.java.opts	
Gateway	ApplicationMaster Memory	yarn.app.mapreduce.am.resource.mb	
Gateway	ApplicationMaster Java Opts Base	yarn.app.mapreduce.am.command-opts	
Gateway	I/O Sort Memory Buffer (MiB)	mapreduce.task.io.sort.mb	

Gateway	Mapreduce Task out	mapreduce.task.out	
Gateway	I/O Sort Factor	mapreduce.task.io.sort.factor	
Gateway	Default Number of Parallel Transfers During Shuffle	mapreduce.reduce.shuffle.parallelcopies	
ResourceManager	Java Heap Size of ResourceManager in Bytes		1024MB

HDFS TESTS

The following tests are used for HDFS benchmark testing.

1. TestDFSIO

The TestDFSIO benchmark is a read and write test for HDFS. It is helpful for tasks such as stress testing HDFS, to discover performance bottlenecks in your network, to shake out the hardware, OS and Hadoop setup of your cluster machines (particularly the NameNode and the DataNodes) and to give you a first impression of how fast your cluster is in terms of I/O.

Scenarios 1: Run the TestDFSIO on 10 Files of 10MB each.

Write 10 Files of 10MB each = total 100MB

```
hadoop jar /usr/hdp/2.4.2.0-258/hadoop-mapreduce/hadoop-mapreduce-client-jobclient-tests.jar TestDFSIO -write -nrFiles 10 -fileSize 10
```

Read 10 Files of 10MB each = total 100MB

```
hadoop jar /usr/hdp/2.4.2.0-258/hadoop-mapreduce/hadoop-mapreduce-client-jobclient-tests.jar TestDFSIO -read -nrFiles 10 -fileSize 10
```

Clean up and remove test data

hadoop jar /usr/hdp/2.4.2.0-258/hadoop-mapreduce/hadoop-mapreduce-client-jobclient-tests.jarTestDFSIO –clean

Scenarios 2: Run the TestDFSIO on 100 Files of 10MB each.

Write 100 Files of 10MB each = total 100MB

hadoop jar /usr/hdp/2.4.2.0-258/hadoop-mapreduce/hadoop-mapreduce-client-jobclient-tests.jarTestDFSIO -write -nrFiles 100 -fileSize 10

Read 100 Files of 1GB each = total 100GB

hadoop jar /usr/hdp/2.4.2.0-258/hadoop-mapreduce/hadoop-mapreduce-client-jobclient-tests.jarTestDFSIO -read -nrFiles 100 -fileSize 10

Clean up and remove test data

hadoop jar /usr/hdp/2.4.2.0-258/hadoop-mapreduce/hadoop-mapreduce-client-jobclient-tests.jarTestDFSIO –clean

Results:

USECASES

- Write 10 Files of 10MB each = total 100MB
- Read 10 Files of 10MB each = total 100MB
- Write 100 Files of 10MB each = total 1GB
- Read 100 Files of 10MB each = total 1GB

Use case	Throughput [mp/sec]	Avg I/O Rate [mb/sec]	IO rate std deviation	Total exec [sec]
Write 10 Files of 10 MB each = total 100MB	71.67739438 335937	72.71614074707 031	8.59193639 0942998	217.23
Read 10 Files of 10MB each = total 100MB	100.0300090 0270081	100.3274765014 6484	5.30274820 593243	153.679

Write 100 Files of 10MB each = total 1GB	74.53753188 342927	78.69841003417 969	26.0812040 25209914	1829.841
Read 100 Files of 10MB each = total 1GB	90.67029831 434849	91.14476776123 047	6.57554218 7824743	1521.729

REMARKS

The TestDFSIO executed successfully on the cluster on both the scenarios, no performance issues were observed during the execution.

2. TERASORT BENCHMARK SUITE

TeraSort is to sort 1TB of data (or any other amount of data you want) as fast as possible. It is a benchmark that combines testing the HDFS and MapReduce layers of a Hadoop cluster.

A full TeraSort benchmark run consists of the following three steps:

1. Generating the input data via TeraGen.
2. Running the actual TeraSort on the input data.
3. Validating the sorted output data via TeraValidate.

To run the Terasort Benchmark test, execute run_terasort.sh



run_terasort.sh

Scenario 1: Run Terasort benchmark test with mapreduce.task.io.sort.mb to 10MB

Scenario 2: Run Terasort benchmark test with mapreduce.task.io.sort.mb to 100MB

Operation	Data	Total exec (ms)
TeraGen	10MB	7294
TeraSort	10MB	7002
TeraValidate	10MB	8576
TeraGen	100MB	8566
TeraSort	100MB	7127
TeraValidate	100MB	10999

3. NAMENODE BENCHMARK (NNBENCH)

NNBench is useful for load testing the NameNode hardware and configuration. It generates a lot of HDFS-related requests with normally very small “payloads” for the sole purpose of putting a high HDFS management stress on the NameNode. The benchmark can simulate requests for creating, reading, renaming and deleting files on HDFS.

Scenarios 1: Run NNBench Create_Write and Open_Read scenarios with 1 file

a. Create_Write with 1 File

```
hadoop jar /usr/hdp/2.4.2.0-258/hadoop-mapreduce/hadoop-mapreduce-client-jobclient-tests.jar nnbench -operation create_write - blockSize 1048576
```

b. Open_Read with 1 File

```
hadoop jar /usr/hdp/2.4.2.0-258/hadoop-mapreduce/hadoop-mapreduce-client-jobclient-tests.jar nnbench -operation open_read - blockSize 1048576
```

Results:

Scenarios	Parameter	TPS	Avg exe	Avg Lat(ms)	Avg Lat (ms)	AL Total I #1	AL Tot	TPS Total	Longest	Lat e	# of exceptions	Remarks
-----------	-----------	-----	------------	--------------------	--------------------	---------------------	-----------	--------------	---------	----------	--------------------	---------

			c (ms)				al #2	l (ms)	Map (ms)	ma ps		
Create_Write with 1 File	nnbench - operation create_write -maps 1 -reduces 1 - blockSize 1048576 - bytesToWrite 0 - numberOfFiles 1 - replicationFactor PerFile 1 - readFileAfterOpen false -baseDir /benchmarks/NN Bench	(Create/Write) 90	22.0	(Create/Write) 18.0	(Close) 4.0	18	4	22	22.0	0	0	real 2m7.150s user 0m4.011s sys 0m0.287s
Open_Read with 1 File	nnbench - operation open_read	(Open/Read) 333	(Open/Read) 3.0	(Open) 3.0	(Read) 0.0	3	0	3	3.0	0	0	real 2m7.177s user 0m4.003s sys 0m0.280s

-maps 1 -reduces 1 -blockSize 1048576 -bytesToWrite 0 -numberOfFiles 1 -replicationFactor PerFile 1 -readFileAfterOpen true -baseDir /benchmarks/NNBench											
---	--	--	--	--	--	--	--	--	--	--	--

Scenarios 2: Run NNBench Create_Write and Open_Read scenarios with 1000 file

c. Create_Write with 1000 File

```
hadoop jar /usr/hdp/2.4.2.0-258/hadoop-mapreduce/hadoop-mapreduce-client-jobclient-tests.jar nnbench -operation create_write -numberOfFiles 1000 -blockSize 1048576
```

d. Open_Read with 1000 File

```
hadoop /usr/hdp/2.4.2.0-258/hadoop-mapreduce/hadoop-mapreduce-client-jobclient-tests.jar nnbench -operation open_read -numberOfFiles 1000 -blockSize 1048576
```

Results:

Scenario	Parameter	TPS	Avg exec (ms)	Avg Lat(ms)	Avg Lat (ms)	AL Total I #1	AL Total #2	TPS Total I (ms)	Longest Map (ms)	Lat e maps	# of exceptions	Remarks
Create_Write with 1000 File	nnbench -operation create_write -maps 1 -reduces 1 -blockSize 1048576 - bytesToWrite 0 - numberOf Files 1000 - replication FactorPerFile 1 - readFileAfterOpen false -baseDir /benchmarks/NNBenchmark	(Create/Write) 978	2.043	(Create/Write) 1.108	(Close) 0.872	1108	872	2043	2043.0	0	0	real 2m10.385s user 0m4.376s sys 0m0.289s
Open_Read with 1000 File	nnbench -operation open_read -maps 1 -reduces 1 -blockSize 1048576 - bytesToWrite 0	(Open/Read) 1968	(Open/Read) 0.508	(Open) 0.453	(Read) 0.0	453	0	508	508.0	0	0	real 2m7.151s user 0m4.011s sys 0m0.236s

- numberOf Files 1000 - replication FactorPerFi le 1 - readFileAft erOpen true -baseDir /benchmar ks/NNBenc h											
---	--	--	--	--	--	--	--	--	--	--	--

Remarks:

The NNBench benchmark executed successfully on the cluster, no read/write exceptions found in both the scenarios.

4. MAP REDUCE TESTS

- MapReduce Benchmark (MRBench)**

MRBench loops a small job a number of times. As such it is a very complimentary benchmark to the “large-scale” TeraSort benchmark suite because MRBench checks whether small job runs are responsive and running efficiently on your cluster. It puts its focus on the MapReduce layer as its impact on the HDFS layer is very limited.

```
hadoop jar /usr/hdp/2.4.2.0-258/hadoop-mapreduce/hadoop-mapreduce-client-jobclient-2.7.1.2.4.2.0-258-tests.jar mrbench -numRuns 50
```

Results:

Parameters	DataLines	Maps	Reduces	Avg (ms)	Remarks
------------	-----------	------	---------	----------	---------

numRuns 50	1	2	1	13568	real	11m20.226s
					user	0m10.888s
					sys	0m0.789s

Remarks:

The MRBench executed successfully with number of runs 50 and average execution 13.56 sec.

CONCLUSION

The benchmark tests executed successfully, we haven't observed any issue.

REFERENCES

- <http://www.michael-noll.com/blog/2011/04/09/benchmarking-and-stress-testing-an-hadoop-cluster-with-terasort-testdfsio-nnbench-mrbench/#namenode-benchmark-nnbench>
- <http://epaulson.github.io/HadoopInternals/benchmarks.html#nnbench>